

CS 2024 FINAL PROJECT USER MANUAL

KIRILL CHERNYSHOV

1. OVERVIEW

This is a project that performs time-benchmarking of several sorting algorithms. The algorithms implemented are:

Algorithm	Average case	Best case	Worst case
Bubblesort	$O(n^2)$	$O(n)$	$O(n^2)$
Insertion sort	$O(n^2)$	$O(n)$	$O(n^2)$
Selection sort	$O(n^2)$	$O(n^2)$	$O(n^2)$
Heapsort	$O(n \log n)$	$O(n)$	$O(n \log n)$
Mergesort	$O(n \log n)$	$O(n \log n)$	$O(n \log n)$
Quicksort	$O(n \log n)$	$O(n^2)$	$O(n \log n)$

The project provides an interface to call the different algorithms on generated data and measure the execution time. The data can be generated with a specified generation rule, or it can be custom data provided by the user.

2. INTERFACE

To use this functionality, first include `SortBenchmark.h`. The algorithms are exposed as static methods, for example `SortBenchmark::bubblesort`. Each of these methods has the following signature:

```
void SortBenchmark::bubblesort(int* array, int length);
```

The two arguments are the array of `int`s to sort and its length. The array is sorted “in-place”: even if the algorithm is not in-place, the array itself is modified, and nothing is returned.

To benchmark individual runs, use `SortBenchmark::runAlgoBenchmark`, which has the following signature:

```
static long long int runAlgoBenchmark(  
    void (* algo)(int*, int),  
    int* array,  
    int length  
);
```

Once again, an array and its length must be provided. The first parameter is for the sorting algorithm, as a pointer. The method returns the time taken in *microseconds* ($1\mu s = 10^{-6}s$). For example, this benchmarks bubblesort:

```

long time = runAlgoBenchmark(
    SortBenchmark::bubblesort,
    int* array,
    int length
);

```

2.1. Mass benchmarking. The core of the class is mass-benchmarking of multiple runs of an algorithm. The class allows to run a specified amount of iterations of a certain algorithm, on arrays with a specified length. This functionality is provided by `SortBenchmark::runAlgoBenchmark`. The array data is generated randomly, but this can be changed.

2.1.1. Iterations and array length. Aside from the algorithm, two more things must be specified: the number of iterations and the array length. The former is how many times to run the algorithm; the total execution time will be benchmarked, which is the sum of times from all the runs. The latter is how big the arrays given to the argument should be. A recommended value that seems to work well on most systems is 100 iterations of 10000-long arrays.

2.1.2. Data generation rules. There are four array generation rules. The first, `SortBenchmark::random_values` is fairly self-explanatory: the data is generated completely randomly with (hopefully) lack of any kind of order. `SortBenchmark::almost_sorted` generates data that is not completely in order, but almost. This is done by generating every element randomly from a small range around the value the element would normally be if the array were sorted. For example, if in a sorted array a value would be 1000, then in an almost-sorted array it would be a random value from 990 to 1010. `SortBenchmark::reverse` simply generates the values in reverse order. This catches out quicksort and bubblesort, which are exceptionally bad at dealing with reverse data. Finally, `SortBenchmark::few_unique` generates an array randomly, but from a small range of values, for example an array of 10000 with numbers ranging from 0 to 99. Note that specifying the generation rule is optional, and will default to `SortBenchmark::random_values`.

2.1.3. Method signature. In general, the method signature is as follows:

```

static long runAlgoBenchmark(
    void (* algo)(int*, int),
    int iterations,
    int arrayLength,
    DataGenRule rule = random_values
);

```

Once again, the return value is the execution time in *microseconds* ($1\mu s = 10^{-6}s$).

2.1.4. *Example.* To run bubblesort on 100 iterations of 10000-long arrays, filled with reverse-sorted data (warning: this will be slow!), use

```
long time = SortBenchmark::runAlgoBenchmark(  
    SortBenchmark::bubblesort ,  
    100 ,  
    10000 ,  
    SortBenchmark::reverse  
);
```

2.2. **Demo.** There is a small, interactive demo provided in main.cpp. Here is an example run:

```
Enter number of iterations: 100  
Enter array length:1000  
Enter the type of data to test on (0 = random, 1 = almost sorted, 2 = reverse, 3 = few unique elements):0  
Bubble sort: 433721 microseconds  
Selection sort: 150809 microseconds  
Insertion sort: 180687 microseconds  
Mergesort: 18610 microseconds  
Heapsort: 31253 microseconds  
Quicksort: 15627 microseconds  
Again? (y/n)
```

The demo is also a good use example of some of the functions.